

Task 2: Values Audit

Task 2: (1) company mission and profile; (2) whose perspective the tool serves; (3) genuine risk vs documentation compliance; (4) who bears the cost if the tool gets it wrong. There are no correct answers; there are shallow answers.

2.1 Company Mission and Profile

Mission Statement

To help financial institutions deploy AI and statistical models responsibly — turning fragmented regulatory expectations across jurisdictions into operational workflows for model inventory, validation, monitoring, escalation, and evidence retention.

The hypothetical company — ClearPath RegTech Pte. Ltd. — is a Singapore-incorporated B2B software provider at the early-growth stage with 25–40 employees (regulatory specialists, data scientists, platform engineers), targeting Tier 1 and Tier 2 financial institutions in Singapore and the US.

ClearPath’s distinctive competence is regulatory translation: reading a jurisdiction’s primary guidance, extracting operative requirements, encoding them as versioned configuration objects, and surfacing the resulting checks within a compliance workflow platform. This is not a data-science firm building proprietary models — it is a regulatory-operations firm building the governance layer that sits around models built by others.

Core Values

- **Accuracy:** Poor monitoring creates false assurance. We surface real signals. If a metric is noisy or a threshold is arbitrary, we say so.
- **Proportionality:** Not every model deserves the same control burden. Risk-tiered, evidence-based governance over uniform checklists.
- **Transparency:** Governance without intelligible evidence becomes box-ticking. Outputs must be readable by risk managers, auditors, and — where required — by affected consumers.
- **Institutional humility:** No monitoring tool can replace human judgment. We are explicit about what the platform does not do.

Revenue and Scale

Pre-Series A. Target Singapore + US pilot clients in Year 1 at SGD 1.5–2 million combined ARR. Total addressable market for AI model risk management tooling across APAC and the US is estimated at USD 1.2 billion by 2027 (Gartner, 2024). ClearPath’s SAM is the subset of large regulated FIs with mature AI programmes and documented governance gaps.

2.2 Whose Perspective Does the Tool Serve?

The tool primarily serves the second-line risk and compliance function: model risk management teams, CCOs, and senior governance committees. These stakeholders have budget authority, mandate for documented oversight, and the power to require remediation of first-line model owners.

However, the tool cannot be designed purely for the paying client’s convenience. There is a genuine three-way tension: the bank (wants efficient deployment and clean sign-off), the regulator (wants robust governance with evidence of substance), and the consumer (wants fair, intelligible, and contestable outcomes when models affect their access to credit). A tool designed solely around the CCO’s desire for a clean dashboard may produce administrative neatness while masking real model risk.

Evidence: The CFPB’s 2023 AI adverse action guidance stated explicitly that a lender cannot cite ‘proprietary algorithm’ as a reason for credit denial. Banks seeking to minimise explainability investment (cost reduction) are in direct tension with this regulatory expectation (consumer protection). The tool’s local SHAP module resolves this operationally — but the underlying business pressure remains.

2.3 Genuine Risk Measurement versus Documentation Compliance

Feature included because it improves risk detection (not required): PSI monitoring

PSI monitoring with a six-window rolling time series is included even though neither OCC 2026-13 nor MAS 2024 prescribes PSI as a specific mandated metric. PSI measures whether the model’s output score distribution has shifted from the development window — a directly observable early-warning signal of model degradation. Using jurisdiction-specific retrain thresholds (0.20 SG, 0.25 US) reflects genuine differences in supervisory conservatism. A documentation-first tool would rely on annual validation cycles; this tool surfaces drift signals between cycles.

Feature included for depth: IV/WoE feature analysis

IV/WoE analysis is included for all input features. While neither OCC 2026-13 nor MAS 2024 mandates IV by name, MAS 2024 (p.14) requires developers to “provide justifications for each feature used.” IV provides a numeric, auditable answer: if Duration has IV = 1.91 (Suspicious — expected artifact of label construction), that is a concrete, documented answer to an examiner’s question.

Feature excluded: uniform approval workflow

The tool does not include an automated approval workflow that routes every model through a fixed sign-off checklist regardless of risk profile. Such a feature would generate voluminous records but would not discriminate between a low-risk, well-validated model and a high-risk, poorly understood one. Instead, the tool implements risk-tiered routing: models are classified by a materiality score, and only those above the high-risk threshold receive independent validation escalation. This reflects OCC 2026-13’s emphasis on proportionality and MAS’s materiality assessment framework.

2.4 Who Bears the Cost if the Tool Gets It Wrong?

Failure Mode	Primary Cost-Bearer	Specific Harm
False negative (missed risk)	Consumers + Bank	Borrowers assessed by degraded/biased model without detection. Bank faces supervisory action and reputational damage on external discovery.
False positive (over-alerting)	Business lines + Model owners	Excessive escalations delay model deployment, increase review costs, erode trust in governance function. Alert fatigue makes real risks harder to surface.
Fairness metric mismeasurement	Consumers (protected groups)	A tool reporting DI = 0.82 (passing) when true DI = 0.73 (failing) conceals discriminatory lending patterns. Harm falls on under-represented borrower groups with no visibility.

Jurisdictional misconfiguration	Compliance + Senior management + Consumers	Applying US Reg B logic where MAS FEAT should apply produces wrong monitoring cadence, wrong escalation pathway, wrong adverse action notice — confident but systematically wrong compliance assertions.
--	--	--

The most serious failure mode is jurisdictional misconfiguration: it produces confident outputs that are systematically wrong, not merely noisy. This is why the jurisdiction registry is version-controlled, each config carries an effective date and primary reference, and the assessment engine logs which config version was active when each report was generated.