Plain-Language Design Summary

DBS Jurisdiction-Aware Credit MRM Platform — US (OCC 2026-13) × SG (MAS 2024/2025 AI MRM)

What the tool does

This tool helps a bank like DBS — which operates credit-scoring AI in both Singapore and the United States — check whether its models are being governed correctly in each country. It runs a set of standard checks on the model's outputs and compares the results against each regulator's specific rules. The key insight is that the same model, assessed against the same data, can be passing in one jurisdiction and failing in another, because the two regulators have made genuinely different choices about what good governance means.

Singapore requires that the bank explain each credit decision to the customer affected by it. The United States does not yet require this, though it does require written reasons for any credit denial. Singapore's retraining threshold (PSI ≥ 0.20) is more conservative than the US threshold (PSI ≥ 0.25) — a gap that creates a 'divergence zone' where the same model is simultaneously non-compliant in Singapore and only under monitoring in the US. These are not technical quirks — they reflect real political choices about consumer protection and institutional burden.

Why a spreadsheet cannot do this

A spreadsheet stores numbers. What it cannot do is apply different thresholds automatically depending on which country's rules are active, produce adverse action notices in the correct format for each jurisdiction, detect when a rule has changed and a running model needs re-assessment, or surface in real time when the same model is simultaneously compliant in one country and non-compliant in another. The tool does all of these through a versioned rule configuration system updated without touching the assessment code. This separation of regulatory logic from assessment logic is the core architectural choice.

What it does not do, and why

The tool does not cover macroeconomic stress testing, Basel III capital model validation, GenAI governance under the US regime, or real-time scoring. These omissions are deliberate. Macroeconomic stress and Basel III are separate regulatory tracks with different supervisory channels. GenAI is excluded from the US path because OCC 2026-13 explicitly excludes it — to include it would misrepresent the current US regulatory position. Stated simply: a governance tool that overstates its coverage is worse than one that clearly names its limits. A governance tool that overstates its coverage is worse than one that clearly states what it does not cover.

What I would have added with more time or real data

On the data side: access to DBS internal data: Singapore Credit Bureau (CBS) scores, DBS internal behaviour scores, and MAS macro stress scenarios from the Financial Stability Review. With more development time: a live regulatory feed for automatic notification when MAS or OCC issues new documents; a full Veritas Toolkit 2.0 FEAT assessment module (the open-source Veritas Toolkit v2.0: https://github.com/veritas-toolkit/); and stress-testing of adverse action notices against CFPB complaint data to validate legal defensibility. These are honest acknowledgements of the gap between a classroom prototype and a production-grade compliance platform — not flaws to hide.